
Conditional Generation of French Electoral Manifestos via QLoRA Fine-Tuning of Llama-3.2-1B on the Archelec Corpus

Félix de Champs
NLP Project Report
felix.dechampsdesaintleger@ensae.fr

Abstract

This report presents a project aimed at generating French electoral manifestos (*professions de foi*) conditioned on a candidate’s political affiliation. The Archelec corpus, composed of digitised campaign documents from the French legislative elections of 1981, 1988, and 1993, provides the training data. An exploratory data analysis confirms that partisan membership produces measurable differences in vocabulary, verbosity, and rhetorical structure, which motivates the use of a conditional language model. We fine-tune meta-llama/Llama-3.2-1B-Instruct using QLoRA four-bit quantisation with the Unsloth framework on a Tesla T4 GPU. Evaluation combines intrinsic metrics – perplexity, ROUGE, and BERTScore – with an extrinsic political distinctiveness test based on TF-IDF logistic regression. The fine-tuned model achieves a perplexity of 4.66 against 18.24 for the base model, a BERTScore F1 of 0.710 versus 0.631, and substantially higher ROUGE scores across all variants. However, the political distinctiveness accuracy of 19% reveals that the generated texts do not yet reliably encode party-specific rhetoric, pointing to limitations inherent in a 1B-parameter model and a small, imbalanced corpus.

1 Introduction

The automatic generation of political discourse is a challenging task at the intersection of natural language processing (NLP), political science, and historical archival research. Unlike general-purpose text generation, political manifestos exhibit strongly constrained rhetorical structures: they must simultaneously convey ideology, appeal to a specific electorate, and respect the conventions of a particular historical moment.

The Archelec corpus aggregates scanned and OCR-processed *professions de foi* from three waves of French legislative elections (1981, 1988, 1993). Each document is a pamphlet produced by a candidate presenting their programme and political identity to the voters of their constituency. The corpus includes rich metadata: candidate name, declared profession, supporting party or coalition, electoral district, election round, and date.

The central problem is conditional text generation: given a set of metadata attributes – particularly the political family of the candidate – can a language model generate a plausible and ideologically consistent manifesto? This formulation is relevant for political historians and corpus linguists studying the evolution of French political rhetoric, and constitutes a concrete NLP challenge in low-resource, domain-specific generation.

The project follows a standard machine learning pipeline. Section 2 describes the corpus and the exploratory data analysis. Section 3 reviews related work. Section 4 motivates the model and training

choices. Section 5 details the experimental setup. Section 6 presents and discusses the results. Section 7 concludes.

2 Data: the Archelec Corpus

2.1 Corpus description and preprocessing

Archelec is a digitised archive of French electoral propaganda hosted by the CEVIPOF (Centre de recherches politiques de Sciences Po). The subset used covers the legislative elections of 1981, 1988, and 1993 and consists exclusively of *professions de foi* files. Each file contains the OCR-extracted text of one manifesto, matched via a shared identifier to a metadata CSV that records candidate-level attributes.

The texts are noisy OCR outputs requiring systematic cleaning. A dedicated pipeline was applied: words broken by end-of-line hyphens are reattached by matching the pattern `[word]-\s+[word]`; isolated capital letters produced by OCR failures (e.g., *V AINCRES* for *VAINCRES*) are merged; watermarks referencing the CEVIPOF provenance and legal printing notices are removed line by line; sequences of three or more consecutive newlines and multiple whitespace characters are normalised. This cleaning function is shared between the EDA and fine-tuning notebooks, ensuring reproducibility.

Figure 1 shows the volume and length distribution of the corpus after cleaning. The three election waves contain approximately 3 100, 3 500, and 5 800 documents respectively, for a total of around 12 400 manifestos. The word-count distribution is right-skewed with a median of 534 words. A secondary mode around 1 200 words is visible, corresponding to longer two-sided pamphlets.




Figure 1: Left: number of documents per election year. Right: distribution of word counts per manifesto after OCR cleaning (median = 534 words). The bimodal shape reflects the coexistence of single- and double-sided pamphlets in the archive.

2.2 Political family labelling and class distribution

Each document was annotated with a political family label derived from the *titulaire-soutien* field. A rule-based normalisation function mapped heterogeneous party affiliations to seven categories: *Communiste et Extreme Gauche*, *Socialiste*, *Droite*, *Extreme Droite*, *Ecologiste*, *Centre*, and *Regionaliste*. Documents labelled *Sans Etiquette* or belonging to unclassifiable micro-parties were excluded from training and evaluation to preserve the integrity of the conditioning signal.

The resulting class distribution (Figure 2) is markedly imbalanced. The *Communiste et Extreme Gauche* category holds approximately 2 350 documents, more than twice the second largest class. This dominance reflects the PCF’s systematic presence in virtually every constituency across all three elections covered. At the other extreme, the *Regionaliste* category contains fewer than fifty documents, a consequence of the geographic concentration of regionalist movements. This imbalance is a structural property of the corpus that directly affects the fine-tuning dynamics.




Figure 2: Number of documents per political family in the labelled corpus. The *Communiste et Extreme Gauche* category is more than twice as large as the next class, a structural imbalance that directly affects training.

2.3 Exploratory analysis: verbosity and lexical specificity

Two complementary analyses were conducted to validate that the political family label carries genuine linguistic information – a prerequisite for conditional generation to be a meaningful task.

The first analysis concerns verbosity. Figure 3 shows the word-count distribution per political family as box plots. The *Communiste et Extreme Gauche* family exhibits a median of approximately 760 words and a large interquartile range, substantially above all other families. The *Extreme Droite* family also shows elevated verbosity with high variance, while *Ecologiste* documents are comparatively concise. These differences suggest that political affiliation encodes stylistic information about expected document length, which a conditional model should be able to reproduce.




Figure 3: Word-count distribution per political family. The *Communiste et Extreme Gauche* family exhibits the highest median and widest spread, indicating systematically longer and more argumentative documents.

The second analysis concerns lexical specificity. Figure 4 shows the ten highest-scoring terms per family by mean TF-IDF, computed after French lemmatisation with SpaCy and removal of a custom stopword list tuned to eliminate generic electoral vocabulary. The profiles are unambiguous: communist texts are dominated by *gauche*, *travailleur*, *parti_communiste*, and *salaire*; socialist texts foreground *parti_socialiste*, *francois_mitterrand*, and *changement*; right-wing texts pivot around *union*, *general*, and *rpr*; far-right texts are saturated by *front_national*, *immigre*, and *immigration*. These lexical profiles are sufficiently distinctive to support linear classification, confirming that the political label is a meaningful conditioning variable.




Figure 4: Top-10 terms by mean TF-IDF score for four political families after lemmatisation and stopword removal. Each family presents a distinct lexical profile, validating the use of political family as a conditioning variable for text generation.

3 Related Work

3.1 Political text analysis and generation

Computational analysis of political texts has a long history in NLP. Early work focused on ideological positioning using bag-of-words or scaling methods (1). More recent approaches apply pre-trained language models to stance detection, framing analysis, and legislative text classification (2). On the generation side, (11) demonstrated that GPT-2 could be steered toward specific topical or sentiment attributes via reinforcement learning from human feedback. In the French political context, (3) explored automatic party-specific discourse generation, confirming that partisan conditioning alters both lexical and syntactic outputs. However, most published work relies on large proprietary models and modern corpora; generating historically grounded political text from a small, noisy archive remains underexplored.

3.2 Parameter-efficient fine-tuning

The dominant approach to adapting large language models (LLMs) with limited data is parameter-efficient fine-tuning (PEFT). LoRA (4) injects trainable low-rank decomposition matrices into attention and feed-forward layers while keeping base weights frozen, drastically reducing the number of trainable parameters. QLoRA (5) extends this by loading the base model in four-bit NormalFloat quantisation and storing the adapter in 16-bit precision, enabling fine-tuning on a single consumer-grade GPU. This is particularly relevant in academic settings where compute is constrained.

3.3 Supervised instruction fine-tuning

Supervised Fine-Tuning (SFT) with instruction templates (6; 7) reframes generation as the completion of an assistant turn given a structured system–user prompt. For domain-specific generation, this preserves the chat formatting learned during base model pre-training while focusing the model’s distribution on the target domain. The TRL library’s `SFTTrainer` (8) provides a high-level API including sequence packing to improve GPU utilisation on datasets with short sequences.

3.4 Evaluation of generated text

ROUGE (9) measures n-gram overlap between hypothesis and reference; it is fast and interpretable but penalises paraphrase. BERTScore (10) computes soft token-level similarity using contextualised embeddings, making it more robust to lexical variation. Perplexity measures how well a model assigns probability mass to held-out text. For politically conditioned generation, an additional extrinsic dimension is necessary: whether generated texts are lexically distinguishable by party. This is operationalised as the accuracy of a TF-IDF classifier trained on real manifestos and applied to generated ones, following (12).

4 Proposed Model and Justification

4.1 Base model selection

`meta-llama/Llama-3.2-1B-Instruct` was selected for three reasons. First, its one-billion-parameter footprint is compatible with the Tesla T4 GPU (16 GB VRAM) under QLoRA quantisation. Second, the model was pre-trained on a multilingual corpus including French-language data, providing basic French grammatical and lexical competence without the need for language-specific pre-training. Third, as an instruction-tuned model, it responds coherently to system–user–assistant prompts via its native chat template, simplifying conditioning prompt construction. A larger model such as `Mistral-7B` would yield better generation quality, but the memory budget and limited corpus size make a 1B model a reasonable starting point.

4.2 QLoRA fine-tuning configuration

The Unsloth library was used for its optimised CUDA implementation of LoRA. The adapter targets all major projection matrices: query, key, value, and output projections, as well as the gate, up, and down projections of the feed-forward blocks. This broad targeting has been shown to outperform restricting LoRA to attention matrices alone (5). The rank is set to $r = 16$ with $\alpha = 32$ (effective scaling factor of 2). Dropout is set to zero following the recommendation for small adapters on small datasets.

4.3 Prompt engineering

The conditioning information is encoded as a two-turn instruction prompt. The system message instructs the model to assume the role of an expert in French political rhetoric and to produce a historically plausible manifesto. The user message provides the candidate’s metadata: name, profession, supporting party, election round, date, and electoral district. The expected completion is the original manifesto text followed by the end-of-sequence token. This structure mirrors the instruction-tuning format of the base model and keeps the political family label as an explicit, parseable attribute in the *soutien politique* field.

5 Experiments

5.1 Dataset preparation

After OCR cleaning, party normalisation, and exclusion of ambiguous documents, the usable dataset retains several hundred examples across seven political families. Only families with at least ten examples were kept to ensure statistical meaningfulness. The 95th percentile of (prompt, completion)

pair lengths fell below 1024 tokens, so `max_seq_length = 1024` was used throughout. The dataset was split 90%/10% with seed 42.

5.2 Training configuration

Training used `SFTTrainer` from TRL. Batch size was 4 per GPU with 4 gradient accumulation steps (effective batch size 16). The `paged_adamw_32bit` optimiser was used to reduce optimiser-state VRAM. Mixed precision was `fp16` (`bfloat16` is unsupported on the T4). Sequence packing was enabled to concatenate short sequences within a single batch slot, substantially increasing throughput. Unsloth gradient checkpointing further reduced peak VRAM. Full hyperparameters are listed in Appendix B.

5.3 Evaluation protocol

Four metrics were computed on the held-out test set. Perplexity is the exponentiated mean cross-entropy over all non-padding tokens. ROUGE-1, ROUGE-2, and ROUGE-L F1 compare generated text to the original manifesto. BERTScore precision, recall, and F1 use multilingual BERT with `lang="fr"`. Political distinctiveness is evaluated by training a TF-IDF logistic regression classifier (5 000 features, sublinear TF, $C = 1$) on real training-set manifestos and applying it to texts generated for seven synthetic candidate profiles, one per political family.

6 Results and Discussion

6.1 Quantitative results

Table 1 summarises all evaluation metrics.

Table 1: Evaluation metrics on the Archelec test set. For ROUGE and BERTScore, generated texts are compared to original manifestos. Political accuracy is computed on seven synthetic profiles.

Metric	Base	Fine-tuned	Δ
Perplexity ($\downarrow$)	18.24	4.66	−74.4%
ROUGE-1 F1 ($\uparrow$)	0.225	0.329	+46.2%
ROUGE-2 F1 ($\uparrow$)	0.037	0.087	+135.1%
ROUGE-L F1 ($\uparrow$)	0.100	0.142	+42.0%
BERTScore F1 ($\uparrow$)	0.631	0.710	+12.5%
Political accuracy ($\uparrow$)	—	0.190	—

6.2 Perplexity

The reduction from 18.24 to 4.66 is the most substantial result of the experiment. It indicates that the fine-tuned model assigns significantly higher likelihood to held-out manifesto texts, reflecting genuine domain adaptation to the register, syntactic patterns, and vocabulary of French electoral pamphlets of the 1980s–90s. The base model, despite its multilingual pre-training, treats this archaic specialised register as low probability. Figure 5 illustrates the magnitude of this gap.

6.3 ROUGE scores

All three ROUGE variants improve substantially after fine-tuning (Figure 6). ROUGE-2 more than doubles (0.037 to 0.087), suggesting that the model reproduces not only individual word choices but also short phrasal patterns characteristic of the domain. Absolute values remain low, which is expected in open-ended generation where paraphrase and structural variation are natural and not penalised by design.

6.4 BERTScore

Figure 7 shows the per-document BERTScore F1 distribution. The fine-tuned model’s distribution shifts rightward (mean 0.710 vs. 0.631) and concentrates around the mean, with fewer extreme low-




Figure 5: Perplexity on the Archelec test set. The fine-tuned model achieves a fourfold reduction relative to the base model.




Figure 6: ROUGE-1, ROUGE-2, and ROUGE-L F1 scores. ROUGE-2 shows the largest proportional gain, indicating improved phrasal consistency.

scoring documents. This concentration indicates consistent production of domain-appropriate text, whereas the base model occasionally generates responses semantically distant from any historical manifesto.




Figure 7: Per-document BERTScore F1 distributions. The fine-tuned model’s distribution is shifted rightward and narrower, indicating both higher average quality and greater consistency.

6.5 Political distinctiveness

The political distinctiveness accuracy of 19% requires careful interpretation. With seven classes, random performance is approximately 14%, so the improvement is marginal. Figure 8 shows the confusion matrix: the classifier maps nearly all generated texts to *Communiste et Extreme Gauche*, with only isolated correct attributions for *Centre* and *Socialiste* profiles.




Figure 8: Confusion matrix of the political family classifier applied to generated texts. The model collapses most predictions to the dominant training class.

This result reveals the central limitation of the approach: the model generates fluent and domain-appropriate French political text but does not reliably condition its content on the requested political family. Three factors contribute to this failure. The corpus is severely imbalanced: the *Communiste et Extreme Gauche* family alone accounts for over one third of all labelled documents, biasing the model’s distribution toward its lexical patterns. The 1B-parameter model likely lacks the capacity to simultaneously capture the general manifesto register and seven distinct ideological sub-styles. Finally, the political label is transmitted via a single free-text field in the prompt, and a small model may not learn a sufficiently strong association between party name and rhetorical output from a dataset of this size.

6.6 Summary

The intrinsic metrics demonstrate clear domain adaptation: the fine-tuned model produces text measurably more similar to historical French manifestos. This result is meaningful for corpus augmentation or historical text synthesis. The extrinsic test, however, shows that domain fluency and ideological conditioning are separate objectives. Generating text that sounds like a 1988 electoral pamphlet is a simpler task than generating one that specifically reflects the vocabulary of a communist rather than a centrist candidate.

6.7 Qualitative and linguistic error analysis

To complement the standard evaluation metrics, a deeper comparative analysis of the linguistic properties of real and generated texts was conducted. This analysis highlights specific dimensions where the fine-tuned model diverges from the historical ground truth.

First, semantic similarity was evaluated by computing the cosine distance between the embeddings of generated and real texts within each political family. The results indicate that the model reproduces the semantic space of certain families more faithfully than others. As illustrated in the corresponding distribution, the *Ecologiste* and *Extrem Droite* categories exhibit the highest median similarity scores, both exceeding 0.70. Conversely, the *Sans Etiquette* category shows the lowest similarity, which is consistent with its inherent ideological heterogeneity. The model captures highly cohesive and distinctive ideological lexicons better than diffuse ones.




Figure 9: Semantic similarity distributions between real and generated texts by political family. Families with highly distinctive vocabularies achieve higher median similarity.

Secondly, a sentiment analysis exposes a distinct mode collapse in the generated distributions. While real manifestos are predominantly positive, they still exhibit a measurable spread across lower sentiment scores, reflecting the critiques and negative campaigning typical of opposition candidates. The generated texts, however, are almost exclusively concentrated in the highest positive sentiment category. The fine-tuned model fails to reproduce the nuanced negativity of real political discourse, defaulting instead to an overly optimistic and uniform rhetorical stance.




Figure 10: Distribution of sentiment scores for real versus generated texts. The generated texts collapse into a uniformly positive mode, losing the critical variance present in historical documents.

Furthermore, the grounding of the generated texts in real-world entities is severely degraded. An extraction of named entities using a pre-trained French recognition model reveals a drastic reduction in the frequency of locations, persons, and organisations in the generated manifestos. Real texts contain on average over twenty organisational references and ten distinct locations per document, whereas generated texts produce fewer than six for each category. This suggests that while the model successfully mimics the syntactic structure of a manifesto, it struggles to populate it with the specific local references, rival names, and institutional acronyms that characterise authentic electoral pamphlets. The resulting text is fluent but geographically and historically generic.




Figure 11: Average frequency of named entities per document. Generated texts systematically under-produce specific references to locations, persons, and organisations.

Finally, an analysis of stylistic metrics provides contrasting results regarding the model’s mimicry capabilities. The distribution of the Flesch reading ease score for generated texts almost perfectly overlaps with that of the real corpus. The model has entirely internalised the syntactic complexity, sentence length, and syllable distribution of French political prose from this era. However, the Type-Token Ratio distributions diverge significantly. The generated texts display a much higher lexical richness than the real manifestos. While language models often suffer from reduced vocabulary diversity, this inversion likely stems from the model failing to reproduce the highly repetitive slogans and redundant phrasing that artificially lower the Type-Token Ratio in authentic campaign materials.




Figure 12: Distributions of lexical richness and reading ease. The model perfectly captures the historical readability level but generates texts with anomalously high vocabulary variation.

7 Conclusion

This project demonstrates that QLoRA fine-tuning of a 1B-parameter instruction model on a small, domain-specific corpus is feasible with constrained hardware and produces measurable improvements in domain adaptation. Perplexity drops fourfold and BERTScore F1 improves by twelve percentage points, confirming that the model internalises the register and vocabulary of French electoral manifestos of the 1980s–90s.

The central limitation is the gap between domain fluency and political distinctiveness. Generated texts are credible as manifestos in general but do not consistently encode each party’s ideological

specificity. Addressing this gap in future work could involve scaling to a larger base model (3B or 7B parameters); applying stratified oversampling or data augmentation to correct class imbalance; incorporating retrieval-augmented generation that supplies in-context examples of the target political family at inference time; or using discriminator-guided decoding such as FUDGE (13) to steer generation toward ideologically distinctive tokens.

From a data perspective, extending the corpus to additional election cycles and local elections would increase coverage and reduce imbalance. Investing in post-correction of the noisiest OCR documents would likely yield quality improvements at lower computational cost than scaling the model.

Despite these limitations, this project establishes a complete, reproducible NLP pipeline from raw archival data to quantitative model evaluation, and provides a documented baseline for conditional generation from the Archelec corpus.

References

References

- [1] Laver, M., Benoit, K., and Garry, J. (2003). Extracting policy positions from political texts using words as data. *American Political Science Review*, 97(2), 311–331.
- [2] Klammer, C., Rehbein, I., and Recasens, M. (2022). Framing and agenda-setting in Russian news: A computational analysis of manipulative rhetorical techniques. In *Proceedings of EMNLP 2022*.
- [3] Barbe, T. and Jacquet, G. (2022). Generating partisan political speeches with GPT-2: A French case study. In *Proceedings of the LREC-Workshop on Language Technology for Political and Administrative Documents*.
- [4] Hu, E., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., and Chen, W. (2022). LoRA: Low-Rank Adaptation of Large Language Models. In *Proceedings of ICLR 2022*.
- [5] Dettmers, T., Pagnoni, A., Holtzman, A., and Zettlemoyer, L. (2023). QLoRA: Efficient finetuning of quantized LLMs. In *NeurIPS 2023*.
- [6] Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., et al. (2022). Training language models to follow instructions with human feedback. In *NeurIPS 2022*.
- [7] Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., and Hashimoto, T. B. (2023). Stanford Alpaca: An instruction-following Llama model. https://github.com/tatsu-lab/stanford_alpaca.
- [8] von Werra, L., Belkada, Y., Tunstall, L., Beeching, E., Thrush, T., Lambert, N., and Huang, S. (2022). TRL: Transformer reinforcement learning. <https://github.com/huggingface/trl>.
- [9] Lin, C.-Y. (2004). ROUGE: A package for automatic evaluation of summaries. In *ACL Workshop on Text Summarization Branches Out*.
- [10] Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., and Artzi, Y. (2019). BERTScore: Evaluating text generation with BERT. In *ICLR 2020*.
- [11] Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. (2019). Fine-tuning language models from human preferences. *arXiv:1909.08593*.
- [12] Prabhunoye, S., Tsvetkov, Y., Salakhutdinov, R., and Black, A. W. (2018). Style transfer through back-translation. In *Proceedings of ACL 2018*.
- [13] Yang, K. and Klein, D. (2021). FUDGE: Controlled text generation with future discriminators. In *Proceedings of NAACL 2021*.

A Additional EDA Figures

This appendix presents four supplementary figures produced in `eda_analysis.ipynb` that complement the main text without being essential to the central argument.

Metadata completeness. Figure 13 shows the percentage of missing values per metadata field. The *profession* field is absent in approximately 51% of records, making it the least reliable conditioning variable. The *soutien* field is missing in about 24% of cases; those documents were excluded from training. Name fields are missing in roughly 6% of records, handled by a placeholder. The *departement* field is almost fully complete and is the most reliable geographic conditioning signal. This analysis motivated the decision to include *profession* in the prompt but to treat its absence as an explicit signal ("non mentionne") rather than imputing a value.




Figure 13: Percentage of missing values per metadata field. The *profession* field is absent in about half of all records, while *departement* is nearly complete.

Temporal evolution of the class distribution. Figure 14 shows stacked document counts per political family and election year. The *Ecologiste* family is virtually absent in 1981 and grows dramatically by 1993, reflecting the rise of the French Greens. The *Extrême_Droite* family grows from a marginal presence in 1981 to a substantial category in 1993, coinciding with the electoral breakthrough of the Front National. Left-wing families are present across all three waves with varying balance. This temporal heterogeneity means the training set is not stationary: the model must generalise across ideological and historical variation simultaneously, adding complexity beyond a purely synchronic generation task.




Figure 14: Stacked document counts per political family and election year. The *Ecologiste* and *Extrême_Droite* families grow substantially from 1981 to 1993, reflecting historical changes in French party politics.

Lexical richness. Figure 15 shows the mean Type-Token Ratio (TTR) per political family and its temporal evolution. The *Communiste_et_Extreme_Gauche* family scores the lowest TTR (approximately 0.58), consistent with a formulaic, slogan-heavy style built around a fixed ideological lexicon. The *Ecologiste* family scores the highest (approximately 0.68), reflecting the vocabulary variety required to address environmental issues spanning multiple policy domains. TTR peaks in 1988 and drops in 1993, suggesting a general homogenisation of electoral discourse as party communication became more professionalised.




Figure 15: Left: mean Type-Token Ratio per political family. Right: temporal evolution of mean TTR, which peaks in 1988 and declines in 1993.

TF-IDF analysis by grammatical class. Figure 16 extends the lexical analysis by separating discriminative terms into four grammatical categories: common nouns, proper nouns, verbs, and adjectives, using SpaCy POS tagging. Common nouns highlight thematic focus (*travailleur* and *salaire* for the communist left; *insecurite* and *immigration* for the far right). Proper nouns reflect endorsement strategies: the communist family foregrounds *parti_communiste* and *georges_marchais*; the socialist family foregrounds *francois_mitterrand*. Verbs capture rhetorical stance: communist texts use *defendre* and *payer*; right-wing texts *redresser* and *assurer*. Adjectives are particularly diagnostic for the far right, which uses *national*, *francais*, and *populaire* as primary identity markers. These per-class profiles provide a richer linguistic fingerprint than aggregate TF-IDF and could inform the design of more targeted conditioning mechanisms in future work.

Mots spécifiques par Famille et par Classe Grammaticale




Figure 16: Top discriminative terms per political family broken down by grammatical category. Each category exposes a distinct rhetorical dimension of party-specific language use.

B Training Hyperparameters

Table 2: Complete QLoRA fine-tuning configuration.

Parameter	Value
Base model	meta-llama/Llama-3.2-1B-Instruct
Quantisation	4-bit NormalFloat (NF4)
LoRA rank r	16
LoRA α	32
LoRA dropout	0
Target modules	q, k, v, o proj.; gate, up, down proj.
Max sequence length	1024
Batch size (per GPU)	4
Gradient accumulation steps	4
Effective batch size	16
Optimiser	paged_adamw_32bit
Mixed precision	fp16
Sequence packing	Enabled
Gradient checkpointing	Unsloth strategy
Random seed	42
Train/test split	90 % / 10 %

C OCR Cleaning Pipeline

The OCR pipeline applies six operations in sequence. Word reattachment targets hyphenated line breaks. Isolated capitals are merged. CEVIPOF watermarks and special OCR characters are removed by exact string matching. Legal printing lines (containing *imprimerie*, five-digit postal codes, or old-format phone numbers) are deleted line by line. Three or more consecutive newlines are collapsed to two. Multiple spaces are normalised to one. The same function is used in both `eda_analysis.ipynb` and `main_v3.ipynb`.

D Political Family Normalisation

The normalisation function applies a priority-ordered keyword lookup to the raw *titulaire-soutien* string. Explicit identifiers take precedence (e.g., "Front national" maps to *Extrême Droite*; "Parti communiste" maps to *Communiste et Extrême Gauche*). Ecological parties are identified by keywords such as *vert* or *environnement*, with an explicit exception for hunting associations (CPNT), mapped to *Droite*. Regionalist movements are detected via geographic markers. The *Sans_Etiquette* label is retained in the EDA to measure the scope of unlabelled data but excluded from training and evaluation.

E Example of Fine-tuned model output

Example user prompt :

Rédige la profession de foi à partir des caractéristiques suivantes :

- Candidat : Jack Toupet
- Profession : Informaticien
- Soutien politique : Ligue communiste révolutionnaire
- Élection : Élections législatives, Tour 1
- Date : 1993-03-21
- Département : Yonne

Generated manifesto :

"LEGISLATIVES DE MARS 93 - 2ème CIRCONSCRIPTION DE LA JOYEUX - CHATEAU-SALIN

Jack TOUPET

Informaticien Maire de Chateau-salins Conseiller Régional d'Allier Candidats présentés par la Ligue Communiste Révolutionnaire.

Madame, Mademoiselle, Monsieur,

Le dimanche 14 mars prochain vous allez voter pour élire votre député. Vous allez choisir entre un candidat socialiste ou une candidate du Parti Communiste. Mais le vote que vous ferez est important car il décide de l'échec ou de la victoire du programme socialo-communiste qui a été largement approuvé dans notre circonscription depuis plusieurs années. La droite n'est pas en position définitive mais elle s'appuiera sur les élections législatives afin de mettre en place une alliance avec le gouvernement pour imposer le retour au pouvoir des socialistes et casser toute chance de changement. Vous avez pu voir ce qu'il implique ces deux partis dont les hommes sont responsables de tout malheur économique et moral, ni l'un ni l'autre ne peuvent résoudre vos difficultés quotidiennes. Le chômage, la précarité, l'inflation, le sous-développement de nos pays riches, les atteintes aux libertés et au droit de vivre dignement, l'exclusion sociale... Toutes ces problèmes se multiplient et aggravent, nous sommes confrontés à leur finitude. Les dirigeants écologistes proposent de faire confiance aux mêmes politiciens qui ont fait cet état de choses durant plus de vingt ans. Ils disent qu'ils feront autre chose parce qu'ils croient sincèrement à cette idée fausse selon laquelle "l'environnement doit être protégé". Comment pourrions-nous croire dans leurs promesses lorsqu'on voit comment ils agissent? Et quelles seraient leurs propositions? Ils veulent une Europe sociale où tous seront égaux devant

l'état. Il faut dire non! Il y a aussi beaucoup de gens qui pensent comme moi-même qu'une société sans classes sociales est possible. Je suis convaincu que cela peut se réaliser grâce à une véritable transformation sociale et industrielle..."